

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6gk8qdgf4n4g3pfee6491hc0brc4

Enter your authorization code:

 Mounted at /content/drive

1. Unsupervised Learning

```
1 %matplotlib inline
2 import scipy
3 import numpy as np
4 import itertools
5 import matplotlib.pyplot as plt
6 import time
```

1. Generating the data

First, we will generate some data for this problem. Set the number of points $N = 400$, their dimension $D = 2$, and the number of clusters $K = 2$, and generate data from the distribution $p(x|z = k) = \mathcal{N}(\mu_k, \Sigma_k)$. Sample 200 data points for $k = 1$ and 200 for $k = 2$, with

$$\mu_1 = \begin{bmatrix} 0.1 \\ 0.1 \end{bmatrix}, \mu_2 = \begin{bmatrix} 6.0 \\ 0.1 \end{bmatrix} \text{ and } \Sigma_1 = \Sigma_2 = \begin{bmatrix} 10 & 7 \\ 7 & 10 \end{bmatrix}$$

Here, $N = 400$. Since you generated the data, you already know which sample comes from which class. Run the cell in the IPython notebook to generate the data.

```
1 # TODO: Run this cell to generate the data
2 num_samples = 400
3 cov = np.array([[1., .7], [.7, 1.]]) * 10
4 mean_1 = [.1, .1]
5 mean_2 = [6., .1]
6
7 x_class1 = np.random.multivariate_normal(mean_1, cov, num_samples // 2)
8 x_class2 = np.random.multivariate_normal(mean_2, cov, num_samples // 2)
9 xy_class1 = np.column_stack((x_class1, np.zeros(num_samples // 2)))
10 xy_class2 = np.column_stack((x_class2, np.ones(num_samples // 2)))
11 data_full = np.row_stack((xy_class1, xy_class2))
12 np.random.shuffle(data_full)
13 data = data_full[:, :2]
14 labels = data_full[:, 2]
```

Make a scatter plot of the data points showing the true cluster assignment of each point using different color codes and shape (x for first class and circles for second class):

```
1 # TODO: Make a scatterplot for the data points showing the true cluster assignments of each point
2 # plt.plot(...) # first class, x shape
3 # plt.plot(...) # second class, circle shape
4 plt.plot(x_class1[:, 0], x_class1[:, 1], 'x', c='red')
5 plt.plot(x_class2[:, 0], x_class2[:, 1], 'o', c='green')
6 plt.show()
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6gk8qdgf4n4g3pfee6491hc0brc4

2. Implement and Run K-Means algorithm

Now, we assume that the true class labels are not known. Implement the k-means algorithm for this problem. Write two functions:

`km_assignment_step`, and `km_refitting_step` as given in the lecture (Here, `km_` means k-means). Identify the correct arguments, and the order to run them. Initialize the algorithm with

$$\hat{\mu}_1 = \begin{bmatrix} 0.0 \\ 0.0 \end{bmatrix}, \hat{\mu}_2 = \begin{bmatrix} 1.0 \\ 1.0 \end{bmatrix}$$

and run it until convergence. Show the resulting cluster assignments on a scatter plot either using different color codes or shape or both. Also plot the cost vs. the number of iterations. Report your misclassification error.

```

1 def cost(data, R, Mu):
2     N, D = data.shape
3     K = Mu.shape[1]
4     J = 0
5     for k in range(K):
6         J += np.sum(np.dot(np.linalg.norm(data - np.array([Mu[:, k], ] * N), axis=1)**2, R))
7     return J

1 # TODO: K-Means Assignment Step
2 def km_assignment_step(data, Mu):
3     """ Compute K-Means assignment step
4
5     Args:
6         data: a NxK matrix for the data points
7         Mu: a DxK matrix for the cluster means locations
8
9     Returns:
10        R_new: a NxK matrix of responsibilities
11    """
12    N, D = data.shape # Number of datapoints and dimension of datapoint
13    K = Mu.shape[1] # number of clusters
14    r = np.zeros([N, K])
15    #a matrix of NxK dimension for the distances of the the N datapoints to each of the K cluster centers.
16    for k in range(K):
17        r[:, k] = np.linalg.norm(data - np.array([Mu[:, k], ] * N), axis=1)**2
18    arg_min = np.argmin(r, axis=1) # argmax/argmin along dimension 1
19    R_new = np.zeros([N, K]) # Set to zeros/ones with shape (N, K)
20    R_new[np.arange(N), arg_min] = 1 # Assign to 1
21    return R_new

1 # TODO: K-means Refitting Step
2 def km_refitting_step(data, R, Mu):
3     """ Compute K-Means refitting step.
4
5     Args:
6         data: a NxK matrix for the data points
7         R: a NxK matrix of responsibilities
8         Mu: a DxK matrix for the cluster means locations
9
10    Returns:
11        Mu_new: a DxK matrix for the new cluster means locations
12    """
13    N, D = data.shape
14    K = Mu.shape[1]
15    Mu_new = np.zeros((D, K))
16    for k in range(K):
17        index = np.where(R[:, k] == 1)
18        Mu_new[:, k] = np.mean(data[index, :], axis = 1)
19    return Mu_new

1 N, D = data.shape
2 K = 2
3 max_iter = 100
4 class_init = np.random.binomial(1., .5, size=N)
5 R = np.vstack([class_init, 1 - class_init]).T
6
7 Mu = np.zeros([D, K])
8 Mu[:, 1] = 1.
9 R.T.dot(data), np.sum(R, axis=0)
10
11 iteration = []
12 clst = []
13 for it in range(max_iter):
14     R = km_assignment_step(data, Mu)
15     Mu = km_refitting_step(data, R, Mu)
16     c = cost(data, R, Mu)
17     iteration.append(it)
18     clst.append(c)
19     print(it, c)
20
21 class_1 = np.where(R[:, 0])
22 class_2 = np.where(R[:, 1])

```



```
0 37292.58165797178
1 37250.88887891708
2 37167.849589052785
3 37036.73687303193
4 36928.79657380079
5 36844.57439989758
6 36773.98217006367
7 36739.05559204959
8 36726.52159971727
9 36726.52159971727
10 36726.52159971727
11 36726.52159971727
12 36726.52159971727
13 36726.52159971727
14 36726.52159971727
15 36726.52159971727
16 36726.52159971727
17 36726.52159971727
18 36726.52159971727
19 36726.52159971727
20 36726.52159971727
21 36726.52159971727
22 36726.52159971727
23 36726.52159971727
24 36726.52159971727
25 36726.52159971727
26 36726.52159971727
27 36726.52159971727
28 36726.52159971727
29 36726.52159971727
30 36726.52159971727
31 36726.52159971727
32 36726.52159971727
33 36726.52159971727
34 36726.52159971727
35 36726.52159971727
36 36726.52159971727
37 36726.52159971727
38 36726.52159971727
39 36726.52159971727
40 36726.52159971727
41 36726.52159971727
42 36726.52159971727
43 36726.52159971727
44 36726.52159971727
45 36726.52159971727
46 36726.52159971727
47 36726.52159971727
48 36726.52159971727
49 36726.52159971727
50 36726.52159971727
51 36726.52159971727
52 36726.52159971727
53 36726.52159971727
54 36726.52159971727
55 36726.52159971727
56 36726.52159971727
57 36726.52159971727
58 36726.52159971727
59 36726.52159971727
60 36726.52159971727
61 36726.52159971727
62 36726.52159971727
63 36726.52159971727
64 36726.52159971727
65 36726.52159971727
66 36726.52159971727
67 36726.52159971727
68 36726.52159971727
69 36726.52159971727
70 36726.52159971727
71 36726.52159971727
72 36726.52159971727
73 36726.52159971727
74 36726.52159971727
75 36726.52159971727
76 36726.52159971727
77 36726.52159971727
78 36726.52159971727
79 36726.52159971727
80 36726.52159971727
81 36726.52159971727
82 36726.52159971727
83 36726.52159971727
84 36726.52159971727
85 36726.52159971727
86 36726.52159971727
87 36726.52159971727
88 36726.52159971727
89 36726.52159971727
90 36726.52159971727
91 36726.52159971727
92 36726.52159971727
93 36726.52159971727
94 36726.52159971727
95 36726.52159971727
```

```


96 36726.52159971727
97 36726.52159971727
98 36726.52159971727
99 36726.52159971727

```

```

1 # TODO: Make a scatterplot for the data points showing the K-Means cluster assignments of each point
2 # plt.plot(...) # first class, x shape
3 # plt.plot(...) # second class, circle shape
4 plt.plot(data[class_1, 0], data[class_1,1], 'x', color='red')
5 plt.plot(data[class_2, 0], data[class_2,1], 'o', color='green')
6 plt.show()

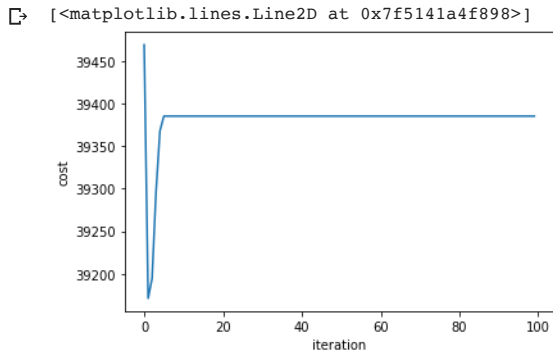
```



```

1 plt.ylabel('cost')
2 plt.xlabel('iteration')
3 plt.plot(iterations, costs)

```



```

1 labels_predict = np.argwhere(R == 1)[: , 1]
2 print("Error rate: ", np.mean(labels != labels_predict))

```

↗

▼ 3. Implement EM algorithm for Gaussian mixtures

Next, implement the EM algorithm for Gaussian mixtures. Write three functions: `log_likelihood`, `gm_e_step`, and `gm_m_step` as given in the lecture. Identify the correct arguments, and the order to run them. Initialize the algorithm with means as in Qs 2.1 k-means initialization, covariances with $\hat{\Sigma}_1 = \hat{\Sigma}_2 = I$, and $\hat{\pi}_1 = \hat{\pi}_2$.

In addition to the update equations in the lecture, for the M (Maximization) step, you also need to use this following equation to update the covariance Σ_k :

$$\hat{\Sigma}_k = \frac{1}{N_k} \sum_{n=1}^N r_k^{(n)} (\mathbf{x}^{(n)} - \hat{\mu}_k)(\mathbf{x}^{(n)} - \hat{\mu}_k)^\top$$

Run the algorithm until convergence and show the resulting cluster assignments on a scatter plot either using different color codes or shape or both. Also plot the log-likelihood vs. the number of iterations. Report your misclassification error.

```

1 def normal_density(x, mu, Sigma):
2     return np.exp(-.5 * np.dot(x - mu, np.linalg.solve(Sigma, x - mu))) \
3         / np.sqrt(np.linalg.det(2 * np.pi * Sigma))

1 def log_likelihood(data, Mu, Sigma, Pi):
2     """ Compute log likelihood on the data given the Gaussian Mixture Parameters.
3
4     Args:
5         data: a NxK matrix for the data points
6         Mu: a DxK matrix for the means of the K Gaussian Mixtures
7         Sigma: a list of size K with each element being DxD covariance matrix
8         Pi: a vector of size K for the mixing coefficients
9
10    Returns:
11        L: a scalar denoting the log likelihood of the data given the Gaussian Mixture
12    """

```

```

12     """
13     # Fill this in:
14     # N, D = ... # Number of datapoints and dimension of datapoint
15     # K = ... # number of mixtures
16     # L, T = 0., 0.
17     # for n in range(N):
18     #     for k in range(K):
19     #         # T += ... # Compute the likelihood from the k-th Gaussian weighted by the mixing coefficients
20     #         L += np.log(T)
21     # return L
22
23 N, D = data.shape # Number of datapoints and dimension of datapoint
24 K = Mu.shape[1] # number of mixtures
25 L, T = 0., 0.
26 for n in range(N):
27     T = 0.
28     for k in range(K):
29         T += Pi[k] * normal_density(data[n], Mu[:, k], Sigma[k]) # Compute the likelihood from the k-th Gaussian weighted
30     L += np.log(T)
31 return L

1 # TODO: Gaussian Mixture Expectation Step
2 def gm_e_step(data, Mu, Sigma, Pi):
3     """ Gaussian Mixture Expectation Step.
4
5     Args:
6         data: a NxN matrix for the data points
7         Mu: a DxK matrix for the means of the K Gaussian Mixtures
8         Sigma: a list of size K with each element being DxD covariance matrix
9         Pi: a vector of size K for the mixing coefficients
10
11     Returns:
12         Gamma: a NxK matrix of responsibilities
13     """
14     # Fill this in:
15     # N, D = ... # Number of datapoints and dimension of datapoint
16     # K = ... # number of mixtures
17     # Gamma = ... # zeros of shape (N,K), matrix of responsibilities
18     # for n in range(N):
19     #     for k in range(K):
20     #         # Gamma[n, k] = ....
21     #         # Gamma[n, :] /= ... # Normalize by sum across second dimension (mixtures)
22     # return Gamma
23 N, D = data.shape # Number of datapoints and dimension of datapoint
24 K = Mu.shape[1] # number of mixtures
25 Gamma = np.zeros((N,K)) # zeros of shape (N,K), matrix of responsibilities
26 for n in range(N):
27     for k in range(K):
28         Gamma[n, k] = Pi[k] * normal_density(data[n], Mu[:,k], Sigma[k])
29     Gamma[n, :] /= np.sum(Gamma[n, :]) # Normalize by sum across second dimension (mixtures)
30 return Gamma

1 # TODO: Gaussian Mixture Maximization Step
2 def gm_m_step(data, Gamma):
3     """ Gaussian Mixture Maximization Step.
4
5     Args:
6         data: a NxN matrix for the data points
7         Gamma: a NxK matrix of responsibilities
8
9     Returns:
10         Mu: a DxK matrix for the means of the K Gaussian Mixtures
11         Sigma: a list of size K with each element being DxD covariance matrix
12         Pi: a vector of size K for the mixing coefficients
13     """
14     N, D = data.shape # Number of datapoints and dimension of datapoint
15     K = Gamma.shape[1] # number of mixtures
16     Nk = np.sum(Gamma, axis=0) # Sum along first axis
17     Mu = np.zeros([D, K])
18     Sigma = np.zeros((K,D,D))
19
20     for k in range(K):
21         for n in range(N):
22             Mu[:,k] += Gamma[n,k] * data[n,:]/Nk[k]
23     for k in range(K):
24         for n in range(N):
25             temp = (data[n,:] - Mu[:,k]).reshape((D,1))
26             Sigma[k] += Gamma[n,k] * np.dot(temp, temp.T)/Nk[k]
27     Pi = Nk / N
28     return Mu, Sigma, Pi

1 # TODO: Run this cell to call the Gaussian Mixture EM algorithm
2 N, D = data.shape
3 K = 2
4 Mu = np.zeros([D, K])

```

```
5 Mu[:, 1] = 1.
6 Sigma = [np.eye(2), np.eye(2)]
7 Pi = np.ones(K) / K
8 Gamma = np.zeros([N, K]) # Gamma is the matrix of responsibilities
9
10 max_iter = 200
11
12 iterations = []
13 costs = []
14 for it in range(max_iter):
15     Gamma = gm_e_step(data, Mu, Sigma, Pi)
16     Mu, Sigma, Pi = gm_m_step(data, Gamma)
17     loglike = log_likelihood(data, Mu, Sigma, Pi)
18     iterations.append(it)
19     costs.append(loglike)
20     print(it, loglike) # This function makes the computation longer, but good for debugging
21
22 class_1 = np.where(Gamma[:, 0] >= .5)
23 class_2 = np.where(Gamma[:, 1] >= .5)
```

↗

0 -2108.569030648646
1 -2103.523067787013
2 -2101.578628940565
3 -2100.39571336378
4 -2099.4236632574143
5 -2098.464055686756
6 -2097.408395697052
7 -2096.1536982077464
8 -2094.5612603236423
9 -2092.519545762133
10 -2090.247433322861
11 -2088.2949298098933
12 -2086.9395731095738
13 -2086.1168033743156
14 -2085.660531138905
15 -2085.421344453549
16 -2085.298761138278
17 -2085.235414098758
18 -2085.2015105004816
19 -2085.1822768473003
20 -2085.1704904611533
21 -2085.1626047237182
22 -2085.1568538546535
23 -2085.1523420208
24 -2085.148604067255
25 -2085.1453911308645
26 -2085.1425645465742
27 -2085.140042670383
28 -2085.137773932699
29 -2085.13572305107
30 -2085.1338639070595
31 -2085.1321758226945
32 -2085.130641581603
33 -2085.1292463507466
34 -2085.1279770703873
35 -2085.1268220897928
36 -2085.1257709341003
37 -2085.1248141428205
38 -2085.1239431490208
39 -2085.1231501829784
40 -2085.122428191501
41 -2085.121770768269
42 -2085.1211720923784
43 -2085.1206268735164
44 -2085.1201303027024
45 -2085.1196780078126
46 -2085.1192660133675
47 -2085.1188907041155
48 -2085.1185487919975
49 -2085.1182372862168
50 -2085.117953466067
51 -2085.117694856276
52 -2085.117459204643
53 -2085.117244461728
54 -2085.117048762411
55 -2085.116870409161
56 -2085.116707856822
57 -2085.1165596988235
58 -2085.1164246546223
59 -2085.1163015583297
60 -2085.116189348338
61 -2085.1160870579333
62 -2085.1159938067362
63 -2085.1159087929304
64 -2085.115831286205
65 -2085.115760621329
66 -2085.115696192309
67 -2085.1156374470756
68 -2085.115583882667
69 -2085.1155350408185
70 -2085.11549050397
71 -2085.115449891633
72 -2085.115412857078
73 -2085.1153790843246
74 -2085.115348285397
75 -2085.115320197835
76 -2085.115294582412
77 -2085.1152712210865
78 -2085.1152499150935
79 -2085.1152304832535
80 -2085.1152127604037
81 -2085.115196595972
82 -2085.1151818526905
83 -2085.115168405407
84 -2085.115156140023
85 -2085.1151449524973
86 -2085.1151347479804
87 -2085.115125439978
88 -2085.1151169496216
89 -2085.115109205
90 -2085.115102140534
91 -2085.115095696413
92 -2085.1150898181154
93 -2085.1150844558997
94 -2085.115079564413
95 -2085.1150751022924

96 -2085.1150710318134
97 -2085.1150673185743
98 -2085.1150639311895
99 -2085.115060841046
100 -2085.115058022042
101 -2085.1150554503643
102 -2085.1150531043013
103 -2085.1150509640465
104 -2085.1150490115356
105 -2085.1150472302825
106 -2085.1150456052655
107 -2085.1150441227674
108 -2085.1150427702855
109 -2085.1150415364123
110 -2085.1150404107357
111 -2085.115039383769
112 -2085.1150384468465
113 -2085.115037592078
114 -2085.1150368122526
115 -2085.1150361007944
116 -2085.11503545171
117 -2085.1150348595315
118 -2085.115034319264
119 -2085.115033826356
120 -2085.1150333766577
121 -2085.115032966381
122 -2085.1150325920626
123 -2085.1150322505555
124 -2085.115031938981
125 -2085.115031654714
126 -2085.115031395365
127 -2085.115031158744
128 -2085.115030942858
129 -2085.115030745898
130 -2085.115030566197
131 -2085.115030402244
132 -2085.11503025266
133 -2085.115030116183
134 -2085.1150299916653
135 -2085.115029878061
136 -2085.1150297744116
137 -2085.1150296798455
138 -2085.115029593564
139 -2085.1150295148445
140 -2085.115029443021
141 -2085.115029377494
142 -2085.115029317707
143 -2085.1150292631587
144 -2085.11502921339
145 -2085.11502916798
146 -2085.115029126552
147 -2085.1150290887545
148 -2085.1150290542655
149 -2085.1150290228
150 -2085.1150289940915
151 -2085.115028967899
152 -2085.115028943999
153 -2085.115028922195
154 -2085.1150289023008
155 -2085.1150288841523
156 -2085.115028867589
157 -2085.1150288524777
158 -2085.1150288386934
159 -2085.1150288261147
160 -2085.1150288146396
161 -2085.115028804169
162 -2085.115028794612
163 -2085.1150287858973
164 -2085.115028777943
165 -2085.1150287706896
166 -2085.115028764068
167 -2085.1150287580294
168 -2085.1150287525174
169 -2085.1150287474907
170 -2085.1150287429027
171 -2085.115028738715
172 -2085.1150287348955
173 -2085.115028731412
174 -2085.1150287282317
175 -2085.115028725332
176 -2085.1150287226865
177 -2085.115028720271
178 -2085.115028718067
179 -2085.115028716056
180 -2085.115028714223
181 -2085.115028712549
182 -2085.115028711023
183 -2085.1150287096307
184 -2085.115028708359
185 -2085.115028707199
186 -2085.115028706141
187 -2085.1150287051746
188 -2085.1150287042956
189 -2085.1150287034893
190 -2085.1150287027563
191 -2085.1150287020887


```

192 -2085.115028701477
193 -2085.115028700922
194 -2085.115028700415
195 -2085.115028699949
196 -2085.1150286995266
197 -2085.1150286991387
198 -2085.1150286987895
199 -2085.115028698466


```

```

1 # TODO: Make a scatterplot for the data points showing the Gaussian Mixture cluster assignments of each point
2
3 plt.scatter(data[class_1][:, 0], data[class_1][:, 1], marker='x', color='red')
4 plt.scatter(data[class_2][:, 0], data[class_2][:, 1], marker='o', color='green')

```

↳ <matplotlib.collections.PathCollection at 0x7f513ea6bc50>




```

1 plt.ylabel('cost')
2 plt.xlabel('iteration')
3 plt.plot(iterations, costs)

```

↳ [<matplotlib.lines.Line2D at 0x7f5141c02a90>]



```

1 labels_predict_EM = np.argwhere(Gamma >= 0.5)[: , 1]
2 print("Error rate: ", np.mean(labels != labels_predict_EM))

```

↳ Error rate: 0.11

```

1 k_means_error_list = []
2 EM_error_list = []
3 k_means_time_list = []
4 EM_time_list = []
5
6 for i in range(5):
7     num_samples = 400
8     cov = np.array([[1., .7], [.7, 1.]]) * 10
9     mean_1 = [.1, .1]
10    mean_2 = [6., .1]
11
12    x_class1 = np.random.multivariate_normal(mean_1, cov, num_samples // 2)
13    x_class2 = np.random.multivariate_normal(mean_2, cov, num_samples // 2)
14    xy_class1 = np.column_stack((x_class1, np.zeros(num_samples // 2)))
15    xy_class2 = np.column_stack((x_class2, np.ones(num_samples // 2)))
16    data_full = np.row_stack([xy_class1, xy_class2])
17    np.random.shuffle(data_full)
18    data = data_full[:, :2]
19    labels = data_full[:, 2]
20
21    N, D = data.shape
22    K = 2
23    max_iter = 200
24    class_init = np.random.binomial(1., .5, size=N)
25    R = np.vstack([class_init, 1 - class_init]).T
26
27    start_time=time.clock()
28    Mu = np.zeros([D, K])
29    Mu[:, 1] = 1.
30    R.T.dot(data), np.sum(R, axis=0)
31    for it in range(max_iter):

```

```

32     R = km_assignment_step(data, Mu)
33     Mu = km_refitting_step(data, R, Mu)
34     c = cost(data, R, Mu)
35     labels_predict = np.argwhere(R == 1)[: , 1]
36     k_means_error_list.append(np.mean(labels != labels_predict))
37     time_use=time.clock()-start_time
38     k_means_time_list.append(time_use)
39
40     start_time = time.clock()
41     Sigma = [np.eye(2), np.eye(2)]
42     Pi = np.ones(K) / K
43     Gamma = np.zeros([N, K]) # Gamma is the matrix of responsibilities
44     for it in range(max_iter):
45         Gamma = gm_e_step(data, Mu, Sigma, Pi)
46         Mu, Sigma, Pi = gm_m_step(data, Gamma)
47         loglike = log_likelihood(data, Mu, Sigma, Pi)
48     labels_predict = np.argwhere(Gamma >= 0.5)[: , 1]
49     EM_error_list.append(np.mean(labels != labels_predict))
50     time_use=time.clock()-start_time
51     EM_time_list.append(time_use)

1 print("BELOW IS THE RESULT FOR FIVE DIFFERENT DATA REALIZATION")
2 print("=====")
3 print("THE K_MEANS REUSLT IS:", k_means_error_list)
4 print("=====")
5 print("THE K_MEANS TIME IS:", k_means_time_list)
6 print("=====")
7 print("THE EM_ERROR REUSLT IS:", EM_error_list)
8 print("=====")
9 print("THE EM_ERROR REUSLT IS:", EM_time_list)

↳ BELOW IS THE RESULT FOR FIVE DIFFERENT DATA REALIZATION
=====
THE K_MEANS REUSLT IS: [0.2325, 0.2675, 0.2675, 0.2675, 0.2375]
=====
THE K_MEANS TIME IS: [0.14491400000000284, 0.12506600000000039, 0.12583499999999503, 0.12067200000001321, 0.12221499999998287]
=====
THE EM_ERROR REUSLT IS: [0.065, 0.0875, 0.1025, 0.1075, 0.0975]
=====
THE EM_ERROR REUSLT IS: [11.5719010000000025, 11.817532999999969, 11.66621299999997, 11.625350999999966, 11.752907999999999]

```

▼ 4. Comment on findings + additional experiments

Comment on the results:

- Compare the performance of k-Means and EM based on the resulting cluster assignments. **Answer:** Compare the performance of k-Means and EM based on resulting cluster assignments, we could see that they are different. Also, notice the misclassification rate for k-Means is significantly larger than the misclassification rate for EM, we could draw the conclusion that EM has a better clustering result than k-Means.
- Compare the performance of k-Means and EM based on their convergence rate. What is the bottleneck for which method? **Answer:** From the result above we could see that for k-Means the convergence time is between 0.1 point to 0.2 point. For EM method, the convergence time is roughly over 10 seconds. Therefore, k-Means converges way faster than EM, and the bottleneck for EM method is EM is more accurate, but takes longer computation time, and for k-Means the bottleneck is it indeed takes relatively shorter time but the prediction is not as accurate as EM's.
- Experiment with 5 different data realizations (generate new data), run your algorithms, and summarize your findings. Does the algorithm performance depend on different realizations of data? **Answer:** From

▼ 2. Reinforcement Learning

There are 3 files:

1. maze.py: defines the MazeEnv class, the simulation environment which the Q-learning agent will interact in.
2. qlearning.py: defines the qlearn function which you will implement, along with several helper functions. Follow the instructions in the file.
3. plotting_utils.py: defines several plotting and visualization utilities. In particular, you will use plot_steps_vs_iters, plot_several_steps_vs_iters, plot_policy_from_q

```

1 # from qlearning import qlearn
2 # from maze import MazeEnv, ProbabilisticMazeEnv
3 # from plotting_utils import plot_steps_vs_iters, plot_several_steps_vs_iters, plot_policy_from_q

1 import numpy as np
2 import matplotlib
3 import matplotlib.pyplot as plt
4 # from qlearning import *
5 # from maze import *
6
7 # UTILITY FUNCTIONS

```

```

8
9
10 color_cycle = ['#377eb8', '#ff7f00', '#a65628',
11                '#f781bf', '#4daf4a', '#984ea3',
12                '#999999', '#e41a1c', '#dede00']
13
14 def plot_steps_vs_iters(steps_vs_iters, block_size=10):
15     num_iters = len(steps_vs_iters)
16     block_size = 10
17     num_blocks = num_iters // block_size
18     smooted_data = np.zeros(shape=(num_blocks, 1))
19     for i in range(num_blocks):
20         lower = i * block_size
21         upper = lower + 9
22         smooted_data[i] = np.mean(steps_vs_iters[lower:upper])
23
24     plt.figure()
25     plt.title("Steps to goal vs episodes")
26     plt.ylabel("Steps to goal")
27     plt.xlabel("Episodes")
28     plt.plot(np.arange(1,num_iters,block_size), smooted_data, color=color_cycle[0])
29
30     return
31
32 def plot_several_steps_vs_iters(steps_vs_iters_list, label_list, block_size=10):
33     smooted_data_list = []
34     for steps_vs_iters in steps_vs_iters_list:
35         num_iters = len(steps_vs_iters)
36         block_size = 10
37         num_blocks = num_iters // block_size
38         smooted_data = np.zeros(shape=(num_blocks, 1))
39         for i in range(num_blocks):
40             lower = i * block_size
41             upper = lower + 9
42             smooted_data[i] = np.mean(steps_vs_iters[lower:upper])
43         smooted_data_list.append(smooted_data)
44
45     plt.figure()
46     plt.title("Steps to goal vs episodes")
47     plt.ylabel("Steps to goal")
48     plt.xlabel("Episodes")
49     index = 0
50     for label, smooted_data in zip(label_list, smooted_data_list):
51         plt.plot(np.arange(1,num_iters,block_size), smooted_data, label=label, color=color_cycle[index])
52         index += 1
53     plt.legend()
54
55     return
56
57
58 # this function sets color values for
59 # Q table cells depending on expected reward value
60 def get_color(value, min_val, max_val):
61
62     switcher={
63         0:'gray',
64         1:'indigo',
65         2:'darkmagenta',
66         3:'orchid',
67         4:'lightpink',
68     }
69
70     step = (max_val-min_val)/5
71     i = 0
72     color='lightpink'
73
74     for limit in np.arange(min_val, max_val, step):
75         if limit <= value < limit+step:
76             color = switcher.get(i)
77             i+=1
78     return color
79
80
81
82 # get first cell out of the start state
83 def get_next_cell(x1,x2,heatmap,policy_table,xlim=9,ylim=9):
84     up_reward=-10000
85     down_reward=-10000
86     left_reward=-10000
87     right_reward=-10000
88
89     if (x1<ylim):
90         if (policy_table[x1-1][x2]!=3):
91             up_reward = heatmap[x1-1][x2]
92     else:
93         up_reward = -1000

```

```

94
95     if (x1>0):
96         if (policy_table[x1+1][x2]!=0):
97             down_reward = heatmap[x1+1][x2]
98     else:
99         down_reward = -1000
100
101     if (x2>0):
102         if (policy_table[x1][x2-1]!=1):
103             left_reward = heatmap[x1][x2-1]
104
105     else:
106         left_reward = -1000
107
108     if (x2<xlim):
109         if (policy_table[x1][x2+1]!=2):
110             right_reward = heatmap[x1][x2+1]
111
112     else:
113         right_reward = -1000
114
115     rewards = np.array([up_reward, down_reward, left_reward, right_reward])
116     idx = np.argmax(rewards)
117     next_cell = [(x1-1,x2), (x1+1,x2), (x1,x2-1), (x1,x2+1)][idx]
118     choice = ['up', 'down', 'left', 'right']
119     #print ('picking ',choice[idx])
120     return next_cell
121
122
123
124
125 # get coordinates of the cells
126 # on the way from the start to goal state
127 def get_path(x1,x2, policy_table):
128     x_coords = [x1]
129     y_coords = [x2]
130     x1_new = x1
131     x2_new = x2
132
133     i=0
134     num_steps = 0
135     total_cells = len(policy_table)*len(policy_table[0])
136     while (policy_table[x1][x2]!='G') and num_steps < total_cells:
137         if (policy_table[x1][x2]==1): # right
138             x2_new=x2+1
139             #print(i, ' - moving right')
140
141         elif (policy_table[x1][x2]==0):
142             x1_new=x1-1
143             #print(i, ' - moving up')
144
145         elif (policy_table[x1][x2]==3):
146             x1_new=x1+1
147             #print(i, ' - moving down')
148
149         elif (policy_table[x1][x2]==2):
150             x2_new=x2-1
151             #print(i, ' - moving left')
152
153         x1 = x1_new
154         x2 = x2_new
155         x_coords.append(x1)
156         y_coords.append(x2)
157         num_steps += 1
158     return x_coords, y_coords
159
160 # plot Q table
161 # optimal path is highlighted and cells colored by their values
162 def plot_table(env, table_data, heatmap, goal_states, start_state, max_val, min_val, x_coords, y_coords):
163     fig = plt.figure(dpi=80)
164     ax = fig.add_subplot(1,1,1)
165     plt.figure(figsize=(10,10))
166
167     width = len(table_data[0])
168     height = len(table_data)
169
170     new_table = []
171
172     for i in range(height):
173         new_row = []
174
175         for j in range(width):
176             if env.map[i][j] == 0:
177                 new_row.append('')
178             else:
179                 digit = table_data[i][j]

```

```

180         if (digit==0):
181             new_row.append('\u2191') # up
182         elif (digit==1):
183             new_row.append('\u2192') # right
184         elif (digit==2):
185             new_row.append('\u2190') # left
186         elif (digit==3):
187             new_row.append('\u2193') # down
188         elif (digit=='G'):
189             new_row.append('G') # goal state
190         elif (digit=='S'):
191             new_row.append('S') # goal state
192         elif (digit==-1):
193             new_row.append('+') # All four directions
194         else:
195             new_row.append('x') # unknown
196
197     new_table.append(new_row)
198
199     table = ax.table(cellText=new_table, loc='center',cellLoc='center')
200
201     table.scale(1,2)
202
203     for i in range(height):
204         new_row = []
205
206         for j in range(width):
207             if new_table[i][j] == ' ':
208                 table[i, j].set_facecolor('black')
209             else:
210                 table[i, j].set_facecolor(get_color(heatmap[i][j],min_val,max_val))
211
212     for goal_state in goal_states:
213         table[(goal_state[0], goal_state[1])].set_facecolor("limegreen")
214     table[(start_state[0], start_state[1])].set_facecolor("yellow")
215     ax.axis('off')
216     table.set_fontsize(16)
217
218     for i in range(len(x_coords)):
219         table[(x_coords[i], y_coords[i])].get_text().set_color('red')
220     plt.show()
221
222
223 # this function takes 3D Q table as an input
224 # and outputs optimal trajectory table (policy table)
225 # and corresponding expected reward values of different cells (heatmap)
226 def get_policy_table(q_hat_3D, start_state, goal_states):
227     policy_table = []
228     heatmap = []
229
230     for i in range(q_hat_3D.shape[0]):
231         row = []
232         heatmap_row = []
233         for j in range(q_hat_3D.shape[1]):
234
235             heatmap_row.append(np.max(q_hat_3D[i,j,:]))
236
237             for goal_state in goal_states:
238                 if (goal_state[0]==i) and (goal_state[1]==j):
239                     row.append('G')
240
241             if (start_state[0]==i) and (start_state[1]==j):
242                 row.append('S')
243             else:
244                 if np.max(q_hat_3D[i,j,:]) == 0:
245                     row.append(-1) # All zeros
246                 else:
247                     row.append(np.argmax(q_hat_3D[i,j,:]))
248             policy_table.append(row)
249             heatmap.append(heatmap_row)
250
251     return policy_table, heatmap
252
253 def plot_policy_from_q(q_hat, env):
254     q_hat_3D = np.reshape(q_hat, (env.m_size, env.m_size, env.num_actions))
255     max_val = q_hat_3D.max()
256     min_val = q_hat_3D.min()
257     start_state = env.get_coords_from_state(env._get_start_state)
258     goal_states = env._get_goal_state
259     goal_states = [env.get_coords_from_state(goal_state) for goal_state in goal_states]
260     policy_table, heatmap = get_policy_table(q_hat_3D, start_state, goal_states)
261     x,y = get_next_cell(start_state[0],start_state[1],heatmap,policy_table)
262     x_coords, y_coords = get_path(x,y,policy_table)
263     plot_table(env, policy_table, heatmap, goal_states, start_state,max_val,min_val, x_coords, y_coords)
264
265     return

```

```

1 import numpy as np
2 import copy
3 import math
4 import random
5
6 ACTION_MEANING = {
7     0: "UP",
8     1: "RIGHT",
9     2: "LEFT",
10    3: "DOWN",
11 }
12
13 SPACE_MEANING = {
14     1: "ROAD",
15     0: "BARRIER",
16    -1: "GOAL",
17 }
18
19
20 class MazeEnv:
21
22     def __init__(self, start=[6,3], goals=[[1, 8]]):
23         """Deterministic Maze Environment"""
24
25         self.m_size = 10
26         self.reward = 10
27         self.num_actions = 4
28         self.num_states = self.m_size * self.m_size
29
30         self.map = np.ones((self.m_size, self.m_size))
31         self.map[3, 4:9] = 0
32         self.map[4:8, 4] = 0
33         self.map[5, 2:4] = 0
34
35         for goal in goals:
36             self.map[goal[0], goal[1]] = -1
37
38         self.start = start
39         self.goals = goals
40         self.obs = self.start
41
42     def step(self, a):
43         """ Perform a action on the environment
44
45         Args:
46             a (int): action integer
47
48         Returns:
49             obs (list): observation list
50             reward (int): reward for such action
51             done (int): whether the goal is reached
52         """
53         done, reward = False, 0.0
54         next_obs = copy.copy(self.obs)
55
56         if a == 0:
57             next_obs[0] = next_obs[0] - 1
58         elif a == 1:
59             next_obs[1] = next_obs[1] + 1
60         elif a == 2:
61             next_obs[1] = next_obs[1] - 1
62         elif a == 3:
63             next_obs[0] = next_obs[0] + 1
64         else:
65             raise Exception("Action is Not Valid")
66
67         if self.is_valid_obs(next_obs):
68             self.obs = next_obs
69
70         if self.map[self.obs[0], self.obs[1]] == -1:
71             reward = self.reward
72             done = True
73
74         state = self.get_state_from_coords(self.obs[0], self.obs[1])
75
76         return state, reward, done
77
78     def is_valid_obs(self, obs):
79         """ Check whether the observation is valid
80
81         Args:
82             obs (list): observation [x, y]
83
84         Returns:
85             is_valid (bool)

```

```

86         """
87
88         if obs[0] >= self.m_size or obs[0] < 0:
89             return False
90
91         if obs[1] >= self.m_size or obs[1] < 0:
92             return False
93
94         if self.map[obs[0], obs[1]] == 0:
95             return False
96
97         return True
98
99     @property
100     def _get_obs(self):
101         """ Get current observation
102         """
103         return self.obs
104
105     @property
106     def _get_state(self):
107         """ Get current observation
108         """
109         return self.get_state_from_coords(self.obs[0], self.obs[1])
110
111     @property
112     def _get_start_state(self):
113         """ Get the start state
114         """
115         return self.get_state_from_coords(self.start[0], self.start[1])
116
117     @property
118     def _get_goal_state(self):
119         """ Get the start state
120         """
121         goals = []
122         for goal in self.goals:
123             goals.append(self.get_state_from_coords(goal[0], goal[1]))
124         return goals
125
126     def reset(self):
127         """ Reset the observation into starting point
128         """
129         self.obs = self.start
130         state = self.get_state_from_coords(self.obs[0], self.obs[1])
131         return state
132
133     def get_state_from_coords(self, row, col):
134         state = row * self.m_size + col
135         return state
136
137     def get_coords_from_state(self, state):
138         row = int(math.floor(state/self.m_size))
139         col = int(state % self.m_size)
140         return row, col
141
142
143 class ProbabilisticMazeEnv(MazeEnv):
144     """ (Q2.3) Hints: you can refer the implementation in MazeEnv
145     """
146
147     def __init__(self, goals=[[2, 8]], p_random=0.05):
148         """ Probabilistic Maze Environment
149
150         Args:
151             goals (list): list of goals coordinates
152             p_random (float): random action rate
153         """
154         self.m_size = 10
155         self.reward = 10
156         self.num_actions = 4
157         self.num_states = self.m_size * self.m_size
158         self.map = np.ones((self.m_size, self.m_size))
159         self.map[3, 4:9] = 0
160         self.map[4:8, 4] = 0
161         self.map[5, 2:4] = 0
162         for goal in goals:
163             self.map[goal[0], goal[1]] = -1
164         self.goals = goals
165         self.start = [6,3]
166         self.obs = self.start
167         self.p_random = p_random
168
169     def step(self, a):
170         done, reward = False, 0.0
171         next_obs = copy.copy(self.obs)

```

```

172
173     rand = random.uniform(0, 1)
174     if rand <= self.p_random:
175         a = np.random.choice(4)
176
177     if a == 0:
178         next_obs[0] = next_obs[0] - 1
179     elif a == 1:
180         next_obs[1] = next_obs[1] + 1
181     elif a == 2:
182         next_obs[1] = next_obs[1] - 1
183     elif a == 3:
184         next_obs[0] = next_obs[0] + 1
185     else:
186         raise Exception("Action is Not Valid")
187
188     if self.is_valid_obs(next_obs):
189         self.obs = next_obs
190
191     if self.map[self.obs[0], self.obs[1]] == -1:
192         reward = self.reward
193         done = True
194
195     state = self.get_state_from_coords(self.obs[0], self.obs[1])
196
197     return state, reward, done
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999

```

```

1 import numpy as np
2 import math
3 import copy
4
5 def qlearn(env, num_iters, alpha, gamma, epsilon, max_steps, use_softmax_policy, init_beta=None, k_exp_sched=None):
6     """ Runs tabular Q learning algorithm for stochastic environment.
7
8     Args:
9         env: instance of environment object
10        num_iters (int): Number of episodes to run Q-learning algorithm
11        alpha (float): The learning rate between [0,1]
12        gamma (float): Discount factor, between [0,1]
13        epsilon (float): Probability in [0,1] that the agent selects a random move instead of
14            selecting greedily from Q value
15        max_steps (int): Maximum number of steps in the environment per episode
16        use_softmax_policy (bool): Whether to use softmax policy (True) or Epsilon-Greedy (False)
17        init_beta (float): If using stochastic policy, sets the initial beta as the parameter for the softmax
18        k_exp_sched (float): If using stochastic policy, sets hyperparameter for exponential schedule
19            on beta
20
21    Returns:
22        q_hat: A Q-value table shaped [num_states, num_actions] for environment with with num_states
23            number of states (e.g. num rows * num columns for grid) and num_actions number of possible
24            actions (e.g. 4 actions up/down/left/right)
25        steps_vs_iters: An array of size num_iters. Each element denotes the number
26            of steps in the environment that the agent took to get to the goal
27            (capped to max_steps)
28
29    """
30    action_space_size = env.num_actions
31    state_space_size = env.num_states
32    q_hat = np.zeros(shape=(state_space_size, action_space_size))
33    steps_vs_iters = np.zeros(num_iters)
34
35    for i in range(num_iters):
36        # TODO: Initialize current state by resetting the environment
37        # curr_state = ...
38        curr_state = env.reset()
39        num_steps = 0
40        done = False
41
42        # TODO: Keep looping while environment isn't done and less than maximum steps
43        # while ...:
44        while (done != True) and (num_steps < max_steps):
45            num_steps += 1
46
47            # Choose an action using policy derived from either softmax Q-value
48            # or epsilon greedy
49            if use_softmax_policy:
50                assert(init_beta is not None)
51                assert(k_exp_sched is not None)
52
53                beta = beta_exp_schedule(init_beta, i, k_exp_sched)
54                action = softmax_policy(q_hat, beta, curr_state, action_space_size)
55            else:
56                # TODO: Epsilon-greedy
57                action = epsilon_greedy(q_hat, epsilon, curr_state, action_space_size)
58            next_state, reward, done = env.step(action)

```



```

58
59     # TODO: Update Q_value
60     if next_state != curr_state:
61         new_value = alpha * (reward + gamma * np.max(q_hat[next_state, :]) - q_hat[curr_state, action])
62         q_hat[curr_state, action] = q_hat[curr_state, action] + new_value
63         curr_state = next_state
64     steps_vs_iters[i] = num_steps
65
66     return q_hat, steps_vs_iters
67
68
69 def epsilon_greedy(q_hat, epsilon, state, action_space_size):
70     """ Chooses a random action with p_rand_move probability,
71     otherwise choose the action with highest Q value for
72     current observation
73
74     Args:
75         q_hat_3D: A Q-value table shaped [num_rows, num_col, num_actions] for
76         grid environment with num_rows rows and num_col columns and num_actions
77         number of possible actions
78         epsilon (float): Probability in [0,1] that the agent selects a random
79         move instead of selecting greedily from Q value
80         obs: A 2-element array with integer element denoting the row and column
81         that the agent is in
82         action_space_size (int): number of possible actions
83
84     Returns:
85         action (int): A number in the range [0, action_space_size-1]
86         denoting the action the agent will take
87     """
88     # TODO: Implement your code here
89     # Hint: Sample from a uniform distribution and check if the sample is below
90     # a certain threshold
91     prob = np.random.uniform()
92     if np.all(q_hat[state, :] == 0) == True:
93         action = np.random.choice(action_space_size)
94     elif (prob <= epsilon):
95         action = np.random.choice(action_space_size)
96     else:
97         action = q_hat[state,:].argmax()
98     return action
99
100
101 def softmax_policy(q_hat, beta, state, action_space_size):
102     """ Choose action using policy derived from Q, using
103     softmax of the Q values divided by the temperature.
104
105     Args:
106         q_hat: A Q-value table shaped [num_rows, num_col, num_actions] for
107         grid environment with num_rows rows and num_col columns
108         beta (float): Parameter for controlling the stochasticity of the action
109         obs: A 2-element array with integer element denoting the row and column
110         that the agent is in
111
112     Returns:
113         action (int): A number in the range [0, action_space_size-1]
114         denoting the action the agent will take
115     """
116     prob = stable_softmax(beta * q_hat)
117     if all(ele == 0 for ele in q_hat[state]):
118         action = np.random.choice(action_space_size)
119     else:
120         action = np.random.choice(action_space_size, p=prob[state,:])
121     return action
122
123
124 def beta_exp_schedule(init_beta, iteration, k=0.1):
125     beta = init_beta * np.exp(k * iteration)
126     return beta
127
128 def stable_softmax(x, axis=1):
129     """ Numerically stable softmax:
130     softmax(x) = e^x / (sum(e^x))
131                 = e^x / (e^max(x) * sum(e^x/e^max(x)))
132
133     Args:
134         x: An N-dimensional array of floats
135         axis: The axis for normalizing over.
136
137     Returns:
138         output: softmax(x) along the specified dimension
139     """
140     max_x = np.max(x, axis, keepdims=True)
141     z = np.exp(x - max_x)
142     output = z / np.sum(z, axis, keepdims=True)
143

```

144 return output

▼ 1. Basic Q Learning experiments


(a) Run your algorithm several times on the given environment. Use the following hyperparameters:

1. Number of episodes = 200
2. Alpha (α) learning rate = 1.0
3. Maximum number of steps per episode = 100. An episode ends when the agent reaches a goal state, or uses the maximum number of steps per episode
4. Gamma (γ) discount factor = 0.9
5. Epsilon (ϵ) for ϵ -greedy = 0.1 (10% of the time)

```
1 # TODO: Fill this in
2 num_iters = 200
3 alpha = 1.0
4 gamma = 0.9
5 epsilon = 0.1
6 max_steps = 100
7 use_softmax_policy = False
8
9 # TODO: Instantiate the MazeEnv environment with default arguments
10 env = MazeEnv()
11
12 # TODO: Run Q-learning:
13 q_hat, steps_vs_iters = qlearn(env, num_iters, alpha, gamma, epsilon, max_steps, use_softmax_policy)
```


Plot the steps to goal vs training iterations (episodes):

```
1 # TODO: Plot the steps vs iterations
2 plot_steps_vs_iters(steps_vs_iters)
```



Visualize the learned greedy policy from the Q values:

```
1 # TODO: plot the policy from the Q value
2 plot_policy_from_q(q_hat, env)
```



<Figure size 720x720 with 0 Axes>

(b) Run your algorithm by passing in a list of 2 goal locations: (1,8) and (5,6). Note: we are using 0-indexing, where (0,0) is top left corner. Report on the results.

```
1 # TODO: Fill this in (same as before)
2 num_iters = 200
3 alpha = 1.0
4 gamma = 0.9
5 epsilon = 0.1
6 max_steps = 100
7 use_softmax_policy = False
8
```

```

8
9 # TODO: Set the goal
10 goal_locs = [[1, 8], [5, 6]]
11 env = MazeEnv(goals= goal_locs)
12
13 # TODO: Run Q-learning:
14 q_hat, steps_vs_iters = qlearn(env, num_iters, alpha, gamma, epsilon, max_steps, use_softmax_policy)

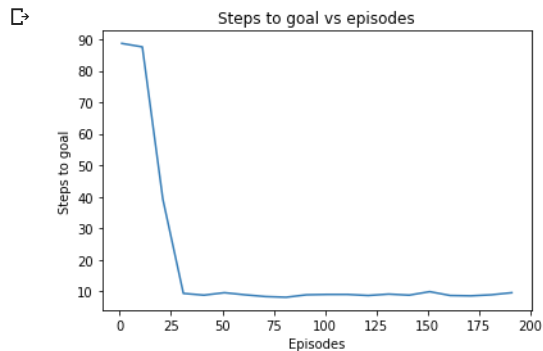
```

Plot the steps to goal vs training iterations (episodes):

```

1 # TODO: Plot the steps vs iterations
2 plot_steps_vs_iters(steps_vs_iters)

```




Plot the steps to goal vs training iterations (episodes):

```

1 # TODO: plot the policy from the Q values
2 plot_policy_from_q(q_hat, env)

```



<Figure size 720x720 with 0 Axes>

▼ 2. Experiment with the exploration strategy, in the original environment


(a) Try different ϵ values in ϵ -greedy exploration: We asked you to use a rate of $\epsilon=10\%$, but try also 50% and 1%. Graph the results (for 3 epsilon values) and discuss the costs and benefits of higher and lower exploration rates.

```

1 # TODO: Fill this in (same as before)
2 num_iters = 200
3 alpha = 1.0
4 gamma = 0.9
5 max_steps = 100
6 use_softmax_policy = False
7
8 # TODO: set the epsilon lists in increasing order:
9 epsilon_list = [0.01, 0.1, 0.5]
10
11 env = MazeEnv()
12
13 steps_vs_iters_list = []
14 for epsilon in epsilon_list:
15     q_hat, steps_vs_iters = qlearn(env, num_iters, alpha, gamma, epsilon, max_steps, use_softmax_policy)
16     steps_vs_iters_list.append(steps_vs_iters)
17
18
19 # TODO: Plot the results
20 label_list = ["epsilon={}".format(eps) for eps in epsilon_list]
21 plot_several_steps_vs_iters(steps_vs_iters_list, label_list)

```

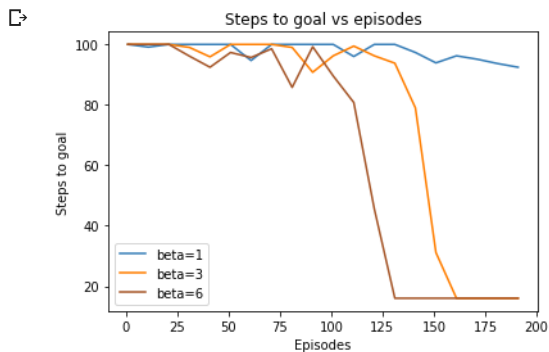
↗



(b) Try exploring with policy derived from **softmax of Q-values** described in the Q learning lecture. Use the values of $\beta \in \{1, 3, 6\}$ for your experiment, keeping β fixed throughout the training.

```
1 # TODO: Fill this in for Static Beta with softmax of Q-values
2 num_iters = 200
3 alpha = 1.0
4 gamma = 0.9
5 epsilon = 0.1
6 max_steps = 100
7
8 # TODO: Set the beta
9 beta_list = [1, 3, 6]
10 use_softmax_policy = True
11 k_exp_schedule = 0.0
12 env = MazeEnv()
13 steps_vs_iters_list = []
14 for beta in beta_list:
15     q_hat, steps_vs_iters = qlearn(env, num_iters, alpha, gamma, epsilon, max_steps, use_softmax_policy, beta, k_exp_schedule)
16     steps_vs_iters_list.append(steps_vs_iters)

1 label_list = ["beta={}".format(beta) for beta in beta_list]
2 # TODO:
3 plot_several_steps_vs_iters(steps_vs_iters_list, label_list)
```



(c) Instead of fixing the $\beta = \beta_0$ to the initial value, we will increase the value of β as the number of episodes t increase:


$$\beta(t) = \beta_0 e^{kt}$$

That is, the β value is fixed for a particular episode. Run the training again for different values of $k \in \{0.05, 0.1, 0.25, 0.5\}$, keeping $\beta_0 = 1.0$. Compare the results obtained with this approach to those obtained with a static β value.

```
1 # TODO: Fill this in for Dynamic Beta
2 num_iters = 200
3 alpha = 1.0
4 gamma = 0.9
5 epsilon = 0.1
6 max_steps = 100
7
8 # TODO: Set the beta
9 beta = 1.0
10 use_softmax_policy = True
11 k_exp_schedule_list = [0.05, 0.1, 0.25, 0.5]
12 env = MazeEnv()
13
14 steps_vs_iters_list = []
15 for k_exp_schedule in k_exp_schedule_list:
16     q_hat, steps_vs_iters = qlearn(env, num_iters, alpha, gamma, epsilon, max_steps, use_softmax_policy, beta, k_exp_schedule)
17     steps_vs_iters_list.append(steps_vs_iters)
```

```
1 # TODO: Plot the steps vs iterations
```

```
2 label_list = ["k={}".format(k_exp_schedule) for k_exp_schedule in k_exp_schedule_list]
3 plot_several_steps_vs_iters(steps_vs_iters_list, label_list)
```



3. Stochastic Environments


(a) Make the environment stochastic (uncertain), such that the agent only has a 95% chance of moving in the chosen direction, and has a 5% chance of moving in some random direction.

```
1 # TODO: Implement ProbabilisticMazeEnv in maze.py
```

(b) Change the learning rule to handle the non-determinism, and experiment with different probability of environment performing random action $p_{rand} \in \{0.05, 0.1, 0.25, 0.5\}$ in this new rule. How does performance vary as the environment becomes more stochastic?

Use the same parameters as in first part, except change the alpha (α) value to be **less than 1**, e.g. 0.5.

```
1 # TODO: Fill this in for Dynamic Beta
2 num_iters = 200
3 alpha = 0.5
4 gamma = 0.9
5 epsilon = 0.1
6 max_steps = 100
7 use_softmax_policy = False
8 beta = 1.0
9
10 # Set the environment probability of random
11 env_p_rand_list = [0.05, 0.1, 0.25, 0.5]
12
13 steps_vs_iters_list = []
14 for env_p_rand in env_p_rand_list:
15     # Instantiate with ProbabilisticMazeEnv
16     env = ProbabilisticMazeEnv(p_random=env_p_rand)
17
18     # Note: We will repeat for several runs of the algorithm to make the result less noisy
19     avg_steps_vs_iters = np.zeros(num_iters)
20     for i in range(10):
21         # q_hat, steps_vs_iters = qlearn(env, num_iters, alpha, gamma, epsilon, max_steps, use_softmax_policy, init_beta=6, k
22         q_hat, steps_vs_iters = qlearn(env, num_iters, alpha, gamma, epsilon, max_steps, use_softmax_policy)
23         avg_steps_vs_iters += steps_vs_iters
24     avg_steps_vs_iters /= 10
25     steps_vs_iters_list.append(avg_steps_vs_iters)
26
27 label_list = ["env_random={}".format(env_p_rand) for env_p_rand in env_p_rand_list]
28 plot_several_steps_vs_iters(steps_vs_iters_list, label_list)
```



3. Did you complete the course evaluation?

ABSOLUTELY YES Thank you so much my dearest professors and TAs.

